

# Firewall Assignment

## Task 1:

We have set-up a network of three VMs using Vbox. Here are the three VMs and the commands run on them:

Host 1:

Public interface (eth0): 172.24.1.1

This is the host which is considered to be on the public internet. This is created using the Bridged networking mode in Vbox. It gets an IP address on the host's network.

Firewall (Host 2):

Public interface (eth0): 172.24.1.2 (example)

Internal Network (eth1): 10.10.10.1

This acts like a router. It has two interfaces: one connected to host through bridge network - eth0, and another on an 'internal network' (another VBox networking mode) - eth1. The firewall port forwards all the connections from it to Host 3 on the internal network. The following iptables rule:

1. ``sudo nano /etc/sysctl.conf`` and write ``net.ipv4.ip_forward=1``
2. `sudo iptables -t nat -A PREROUTING -i eth0 -j DNAT --to-destination 10.10.10.2`. This will DNAT all the traffic from firewall to host 2.
3. `sudo iptables -t nat -A POSTROUTING -o eth1 -d 10.10.10.2 -j SNAT --to-source 10.10.10.1`. This is required so that Host 2 replies back to Firewall instead of to the original ip.

Host 3:

Internal Network (eth0): 10.10.10.2

This is the VM where our server will be running. It is not publicly accessible and can be only accessed through the router/firewall.

We have a working router setup now. Host 1 can reach Host 3 by accessing Firewall's public IP. We run our custom firewall on Host 2 using steps discussed in Task 2.

## Task 2:

We have used NFQUEUE for forwarding all the incoming packets to the host to the firewall code.

This rule

**iptables -I INPUT -j NFQUEUE --queue-num 1**

Is ran to forward all packets from iptables INPUT policy to the NetFilter Queue 1.

In firewall.py, when a packet enters the queue, a callback function is called which validates the packet and performs the required action. First we check if packet has valid headers.

If packet is valid, then rules are iterated over. When a rule succeeds remaining rules are not checked and the required action is performed . This checking happens sequentially. If a packet is being checked, others need to wait for it to finish.

Interactive CLI

Program cli.py is run in parallel which makes a file database.json. When a rule is added via CLI, we add a new entry in the file.


In firewall callback this file is read, thereby reading the latest rules.

## Task 3:

Calculation of PPS : When firewall finishes, this event is logged in log.txt along with the timestamp. When we want the PPS, this file is read & a graph is plotted depicting the PPS over time interval of 1 second.

Testing: iperf tool was used to make the required test-bed and get traffic.

## MAXIMUM PPS over accepting all packets



### Configuration

```
iperf3 --server --port 5201 -f K -V
```

```
iperf3 --client 172.24.1.102 --port 5201 --version4 -f K -V --bandwidth 50M --time 60 --parallel 30
```

### Results:

After testing with various bandwidth and parallel connections : the MAX PPS obtained was 5777. This result and above graph was obtained on accepting all the inputs obtained by iperf.

## Various no of rules matching with traffic equally distributed

### Configuration


The bash files used are submitted along with this assignment to make the configuration.

### Testbed:

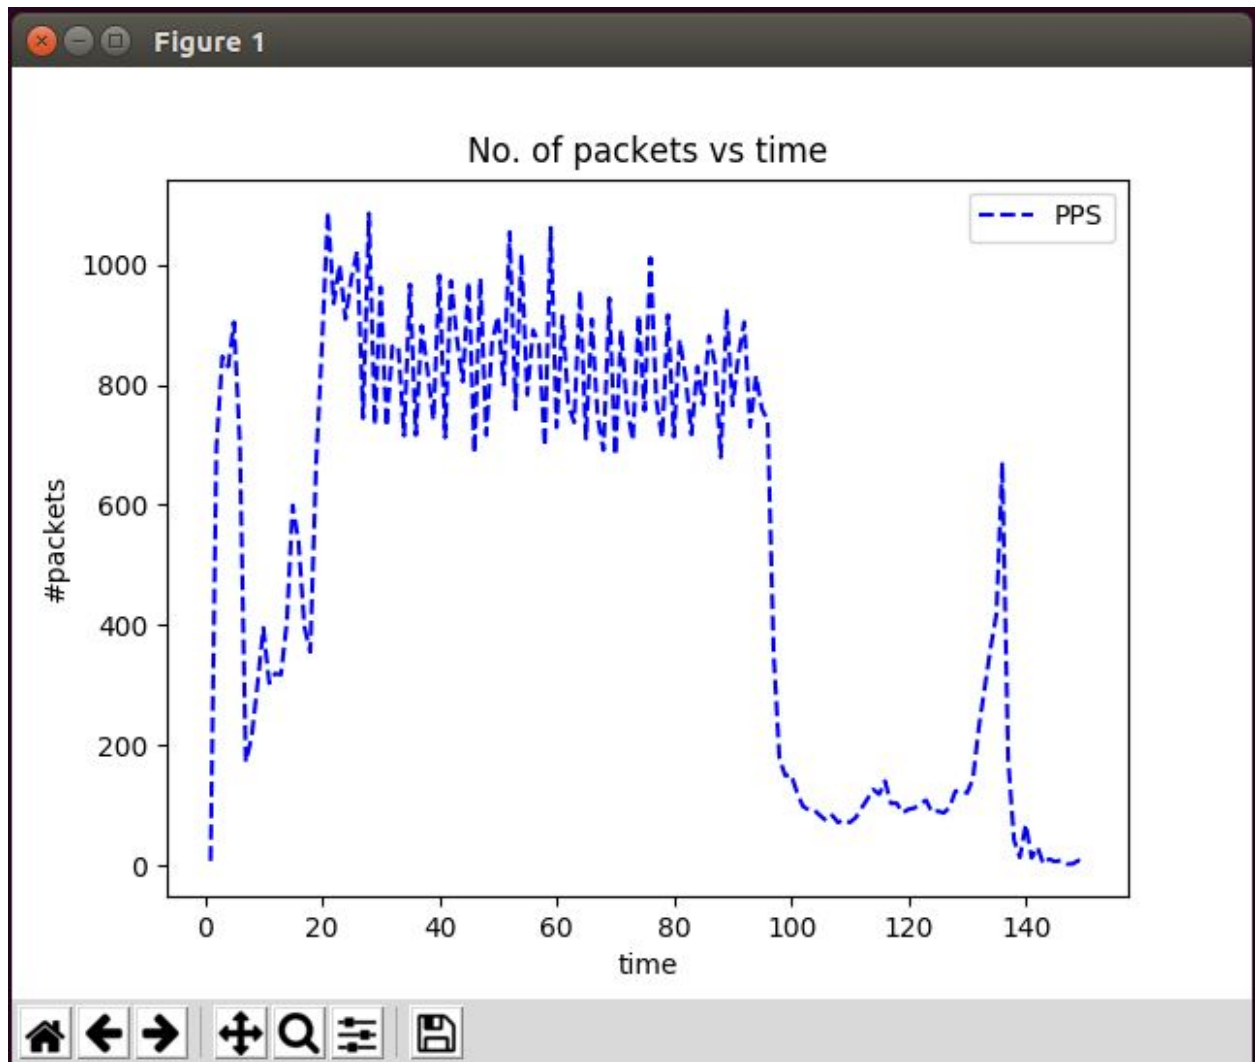
10, 25, 50 TCP rules were made with various destination port. Then those many iperf servers were spawned at the mentioned port. IPerf clients were made to generate traffic on the mentioned ports.


### Results

No. of Rules = 10



No. of Rules = 25





Max PPS :: 541

Observations

With more rules, PPS decreases.

This is because more computation needs to be done per packet. And next packet only goes for processing once previous callback returns

Since results for each experiment were conducted for

```
Terminal File Edit View Search Terminal Tabs Help
root@harsh-Inspiron-5558: /home/harsh/btech/sem-6/Co... x root@hars
Got 13 data points
Maximum PPS 7 at time 1
root:~/# clear
root:~/# show
Chain INPUT (policy DROP)
Rule-No target prot Source-IP
0 ACCEPT all any dpt 9001
1 ACCEPT all any dpt 9002
2 ACCEPT all any dpt 9003
3 ACCEPT all any dpt 9004
4 ACCEPT all any dpt 9005
5 ACCEPT all any dpt 9006
6 ACCEPT all any dpt 9007
7 ACCEPT all any dpt 9008
8 ACCEPT all any dpt 9009
9 ACCEPT all any dpt 9010
10 ACCEPT all any dpt 9011
11 ACCEPT all any dpt 9012
12 ACCEPT all any dpt 9013
13 ACCEPT all any dpt 9014
14 ACCEPT all any dpt 9015
15 ACCEPT all any dpt 9016
16 ACCEPT all any dpt 9017
17 ACCEPT all any dpt 9018
18 ACCEPT all any dpt 9019
19 ACCEPT all any dpt 9020
20 ACCEPT all any dpt 9021
21 ACCEPT all any dpt 9022
22 ACCEPT all any dpt 9023
23 ACCEPT all any dpt 9024
24 ACCEPT all any dpt 9025
25 ACCEPT all any dpt 9026
26 ACCEPT all any dpt 9027
27 ACCEPT all any dpt 9028
28 ACCEPT all any dpt 9029
29 ACCEPT all any dpt 9030
30 ACCEPT all any dpt 9031
31 ACCEPT all any dpt 9032
32 ACCEPT all any dpt 9033
33 ACCEPT all any dpt 9034
34 ACCEPT all any dpt 9035
35 ACCEPT all any dpt 9036
36 ACCEPT all any dpt 9037
37 ACCEPT all any dpt 9038
38 ACCEPT all any dpt 9039
39 ACCEPT all any dpt 9040
40 ACCEPT all any dpt 9041
41 ACCEPT all any dpt 9042
42 ACCEPT all any dpt 9043
43 ACCEPT all any dpt 9044
44 ACCEPT all any dpt 9045
45 ACCEPT all any dpt 9046
46 ACCEPT all any dpt 9047
47 ACCEPT all any dpt 9048
48 ACCEPT all any dpt 9049
49 ACCEPT all any dpt 9050
root:~/# plot
Got 6353 data points
Maximum PPS 541 at time 4
root:~/# plot
```

Rule Table containing 50 rules



```

root@harsh-Inspiron-5558:/home/harsh/btech/sem-6/Computer-Network-Security/Assignments/Assgn-4-Firewall# iptables -L -v --line-number
Chain INPUT (policy ACCEPT 0 packets, 0 bytes)
num  pkts bytes target    prot opt in     out     source destination
1    670K 19G NFQUEUE   all  --  any    any     anywhere anywhere          NFQUEUE num 1

Chain FORWARD (policy ACCEPT 0 packets, 0 bytes)
num  pkts bytes target    prot opt in     out     source destination
1    581 119K DOCKER-USER all  --  any    any     anywhere anywhere
2    581 119K DOCKER-INGRESS all --  any    any     anywhere anywhere
3    581 119K DOCKER-ISOLATION all --  any    any     anywhere anywhere
4      0 0 ACCEPT all  --  any    docker0 anywhere anywhere ctstate RELATED,ESTABLISHED
5      0 0 DOCKER all  --  any    docker0 anywhere anywhere
6      0 0 ACCEPT all  --  docker0 !docker0 anywhere anywhere
7      0 0 ACCEPT all  --  docker0 docker0 anywhere anywhere
8      0 0 ACCEPT all  --  any    docker_gwbridge anywhere anywhere ctstate RELATED,ESTABLISHED
9      0 0 DOCKER all  --  any    docker_gwbridge anywhere anywhere
10     0 0 ACCEPT all  --  docker_gwbridge !docker_gwbridge anywhere anywhere
11     0 0 DROP all  --  docker_gwbridge docker_gwbridge anywhere anywhere

Chain OUTPUT (policy ACCEPT 665K packets, 19G bytes)
num  pkts bytes target    prot opt in     out     source destination

Chain DOCKER (2 references)
num  pkts bytes target    prot opt in     out     source destination

Chain DOCKER-INGRESS (1 references)
num  pkts bytes target    prot opt in     out     source destination
1      0 0 ACCEPT tcp  --  any    any     anywhere anywhere tcp dpt:30000
2      0 0 ACCEPT tcp  --  any    any     anywhere anywhere state RELATED,ESTABLISHED tcp spt:30000
3      0 0 ACCEPT tcp  --  any    any     anywhere anywhere tcp dpt:http
4      0 0 ACCEPT tcp  --  any    any     anywhere anywhere state RELATED,ESTABLISHED tcp spt:http
5    581 119K RETURN all  --  any    any     anywhere anywhere

Chain DOCKER-ISOLATION (1 references)
num  pkts bytes target    prot opt in     out     source destination
1      0 0 DROP all  --  docker_gwbridge docker0 anywhere anywhere
2      0 0 DROP all  --  docker0 docker_gwbridge anywhere anywhere
3    581 119K RETURN all  --  any    any     anywhere anywhere

Chain DOCKER-USER (1 references)
num  pkts bytes target    prot opt in     out     source destination
1    581 119K RETURN all  --  any    any     anywhere anywhere
root@harsh-Inspiron-5558:/home/harsh/btech/sem-6/Computer-Network-Security/Assignments/Assgn-4-Firewall#

```

In the green box, the total bytes of packet is seen that was obtained while testing and getting the results

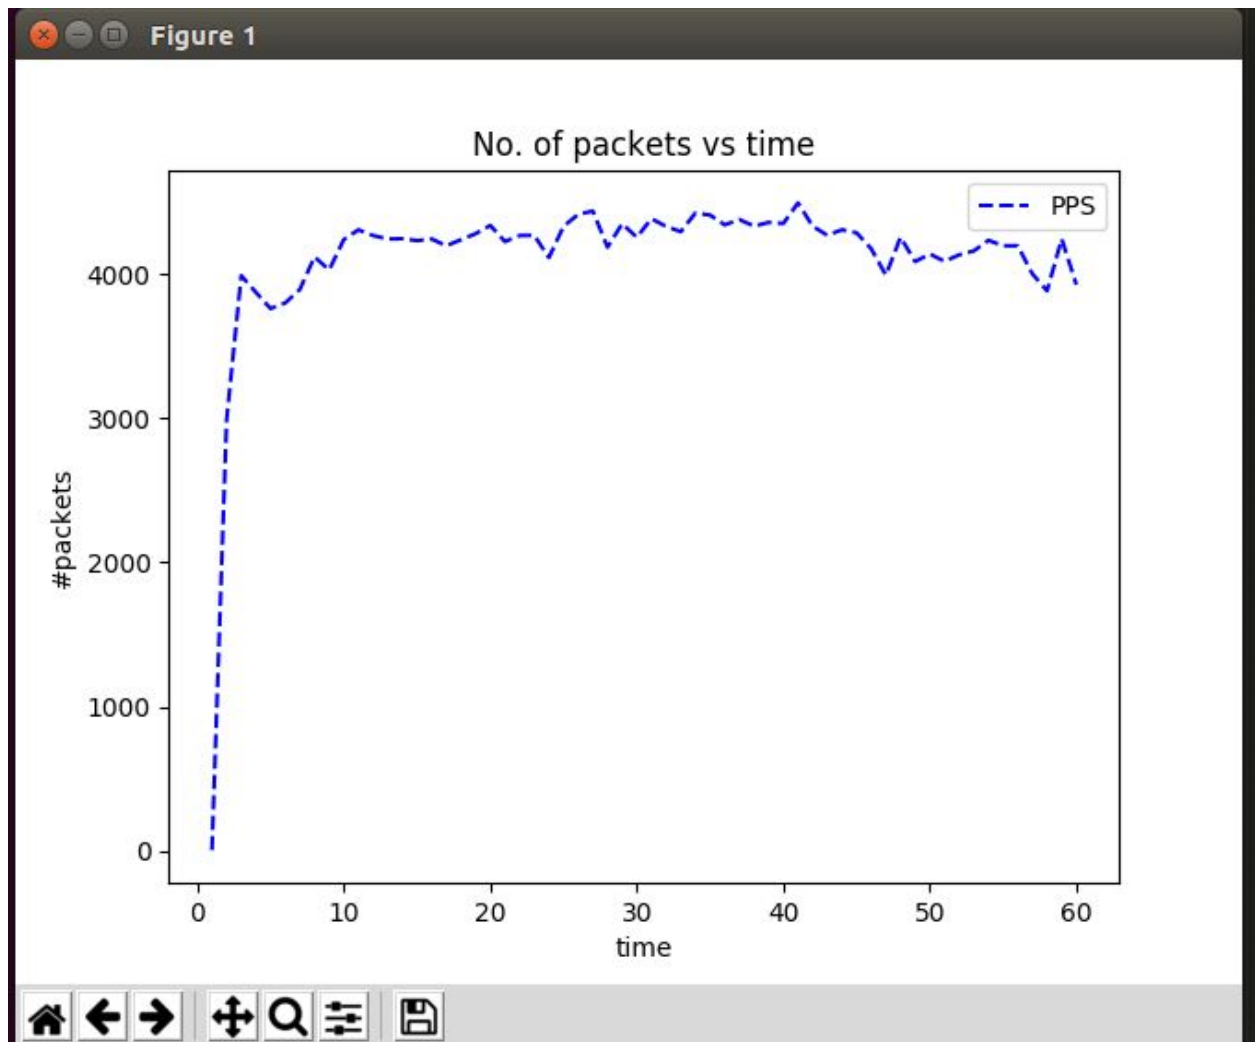


[illegible]

Since callback is sequential, we can see that ports are checked from 9001 to 9020, before an acceptance rule is passed

## Many matching fields

```
RULE :: ADD -dport 5201 -s 127.0.0.1 -p tcp -j accept
iperf3 --server --port 5201 -f K -V
iperf3 --client 127.0.0.1 --port 5201 --version4 -f K -V --bandwidth 50M --time 60 --parallel 30
Result ::
```



Max PPS obtained was 4495.

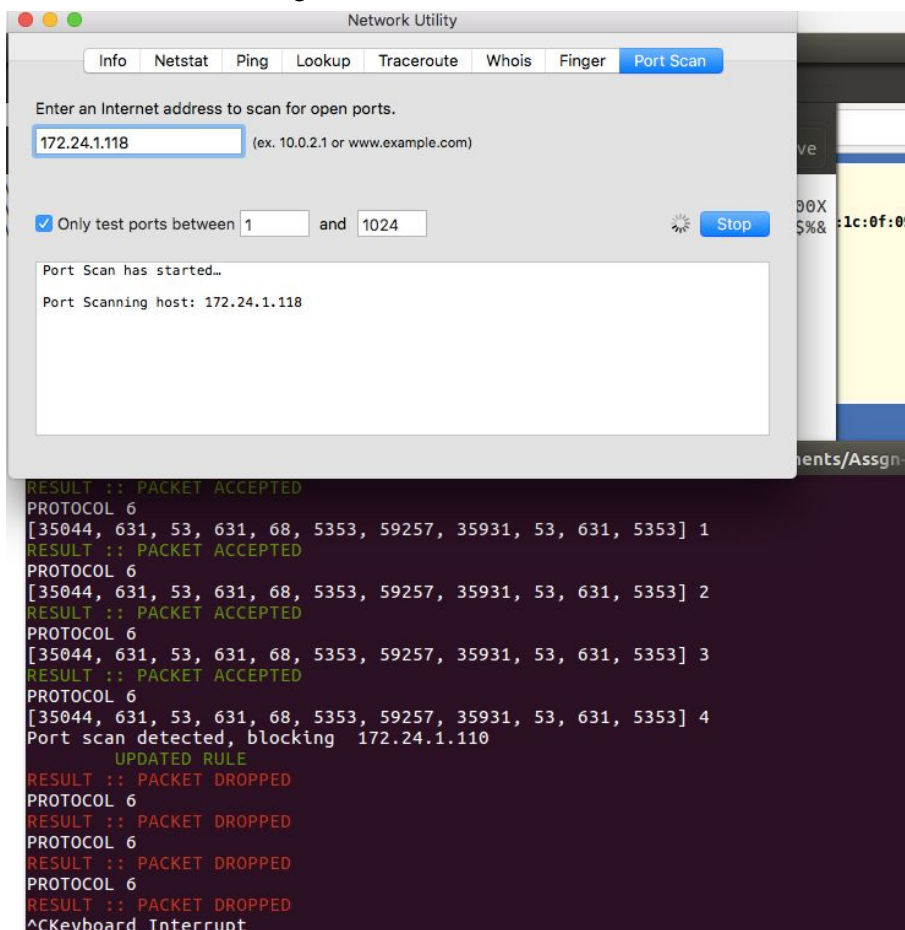
This is less than the rule of Accepting all(5777) because more time is spent per packet.

## Task 4 - Detecting Port Scanning Attacks

The firewall is capable of detecting and blocking port scanning attacks. It does in the following manner. The firewall records all the invalid port scans (ie. the port scans on the ports that were closed at the the time of the request) for every IP address in the last one minute. If the number of invalid port scans exceeds a certain threshold in the last minute, the IP address is blocked. It gets the list of open ports by getting all the open UNIX sockets (similar to `netstat -lntu`), and if we receive a request for port number which is not in that list, we increment the invalid port scans count. To test this:

- Start the firewall, and add the rule “ADD -j ACCEPT” through cli.py.
- Run a port scan using any utility

You should see that it gets blocked and a new DROP rule is added to the firewall.



As can be seen, the ip address of the attacker was blocked after it detected 5 consecutives failed pings.

## Reference

<https://pypi.python.org/pypi/NetfilterQueue> : for using NFQUEUE

<http://iperf.fr/> : for testing & benchmarking